

Using Machine Learning for move sequence visualization and generation in climbing

Rimbot Thomas

EPFL - Semester project 2022, supervised by Martin Jaggi

Abstract—TODO

I. INTRODUCTION

Applying Machine Learning techniques to competitive sport has been an increasing trend in the past few years. We can for example cite the case of car racing or hockey. In this project, we focus on *bouldering*, a form of rock climbing where athletes are tasked with overcoming a small natural or artificial feature (about 4m high), requiring both physical strengths and problem-solving skills.

Previous projects realized by EPFL students (see [1] and [2]) have looked into pose analysis for move sequence detection applied to bouldering. Here, we aim at expanding this research, providing a more tangible visualization pipeline and experimenting with move sequence prediction from holds information.

II. MOVE SEQUENCE DETECTION

The starting point of this project was the work done in [1]. In this work, the authors developed a pipeline for move sequence detection starting from climbing videos. More details can be found in their paper but we refer in this section the general procedure.

A. Pose estimation

Using the Mediapipe [3] library, the video of a person climbing is analyzed, and this person's pose can be extracted. More precisely, for each frame, the visibility and $x - y$ coordinates of the 33 landmarks are acquired, amounting to a total of 99 pose information. The landmarks correspond to specific points on the human body, such as the ankle, knee, elbow etc.

B. Static extremity detection

Once the landmarks information are all recorded, they are aggregated into a dataframe and the points corresponding to common extremities are gathered together. For instance, the left hand extremity is composed of the landmarks corresponding to the left wrist, pinky, index and thumb. By checking the distance between two consecutive positions of an extremity at two different frames, it is possible to evaluate if it was static or not, *i.e.* if it was stable on a hold or in motion. This way, we end up with a cloud of points for all extremities and frames where they are static.



Figure 1: Move sequence detection from [1]

C. Clustering and visualization

From this cloud of points, we can use clustering algorithms, such as the DBSCAN [4] to regroup them and evaluate the position of the holds. Using OpenCV [5], we can draw squares on a still picture of the boulder, representing the detected holds. Moreover, since we have access to the time information through the frame number, we can find the order in which the holds are used, which, combined with the extremity, is exactly the *Move Sequence*.

III. MOVE SEQUENCE VISUALIZATION

Knowing that working directly with the move sequence was possible, we set out to create a tool which would improve the user experience and allow to visualize it in a more tangible way.

A. Skeleton generation

From [1], we can output the move sequence as a sequence of holds coordinates, together with information on the used extremity for each move. This sequence can then be printed on a still image of the boulder, as shown in Figure 1. However, it is hard to visualize and imagine an actual human climbing, following this sequence. Therefore, we created a visualization pipeline, with the goal of generating and drawing a humanoid figure using Mediapipe's features.

More precisely, Mediapipe allows to draw the recorded landmarks, and superimpose them on the video of the person

climbing. These points, linked together with edges, form a humanoid skeleton representing a climber. Playing with this tool, we were able to draw the landmarks directly on the image of the boulder instead of the video.

With this tool in place, the next step was to generate the landmarks information from the move sequence. More precisely, the general form of a move sequence is reported in Table I. The x and y coordinates represent the position of the hold, the limb is the extremity used (hand or foot, left or right), and the indices indicate the order in which these holds are used.

x_1	y_1	$limb_1$
x_2	y_2	$limb_2$
x_3	y_3	$limb_3$
...
x_n	y_n	$limb_n$

Table I: Move sequence template

On the other hand, Mediapipe needs precise information on all 33 landmarks, as explained above, for each frame. So we had to find a way of generating the 99 features for each frame of the desired video duration (which we usually considered to be about 1500 frames), from the move sequence alone. This was done in two steps.

1) Generation of the landmarks sequence for the extremities: The move sequence gives us information on the position of the extremities, in a certain order. From this, we can interpolate the positions between consecutive holds to get a continuous motion of the extremities. For instance, if the four extremities are fixed on holds, and we know that the next move in the sequence involves the left hand, we can interpolate the coordinates of the corresponding landmarks between the current hold's and the next one's. This allows to generate a continuous sequence of landmarks coordinates for the left hand, which can then be used to draw the skeleton with the above-described features.

The interpolation speed is weighted by the distance between two consecutive holds of the same extremity and the average number of desired frames for this move, giving a sense of realism to the motion and allowing to tune its duration.

2) Generation of the landmarks sequence for the rest of the body: Using this procedure, we have a sequence of information for all the extremities landmarks. We are now left to generate the information of the rest of the landmarks, namely the ones corresponding to the rest of the body: the legs, torso, arms and head. In order to do that, we can use the work done in [1] for pose detection. In fact, we have access to a lot of video-recorded landmarks sequence of climbers. From this dataset, we can use the extremities landmarks information as features, and try to predict the body landmarks.

Using scikit-learn [6]'s simple Linear Regression model on a limited number of these videos, we achieve an already very good accuracy of more than 99%, enough for our visualization purposes. From this,



Figure 2: Generated skeleton from a move sequence, using the procedure described in subsection III-A

we now have access to the whole sequence of coordinates for each landmark, for each frame. Setting all the visibilities to 1, we have a complete dataframe of the 99 landmarks information, for the whole video.

Using Mediapipe, we can now draw the skeleton from the landmarks dataframe and visualize the move sequence. The result is a video of a humanoid figure, climbing the boulder according to the specified move sequence. A few frames of such a video are displayed in Figure 2.

The results are definitely satisfactory. We indeed get an idea of how a person would move on the wall, following the given move sequence. However, the use of a linear interpolation implies that the extremities move one after the other, and are fixed when static. Therefore, this can lead to non-physical behaviours, such as limbs stretching, when looking at a dynamic boulder. In general, any problem that would require the climber to lose contact with the wall, or move multiple extremities at the same time, will result in an inaccurate visualization. This is in particular problematic when looking at new-school bouldering, where a lot of jumps and balancing are present. A possible solution to this is an alternative approach, where all the landmarks sequence are directly predicted from the move sequence, instead of relying on linear interpolation for the extremities. This could be the subject of a future work.

B. Selection interface

In order to facilitate the use of move sequences for the skeleton pipeline, we developed a selection interface using OpenCV. The goal was to provide a UI allowing the user to dynamically choose their sequence of movements on a standardized wall, the *Moonboard*. By recording the cursor position and the selected extremity, the move sequence can be generated and displayed on the image. Exported as a dataframe, it can then be processed by the above-described pipeline and visualized as a video with the skeleton.

We report in Figure 3 an example of use of the interface. On the left is displayed a selected move sequence, and on the right a simple holds sequence. The difference between the two lies in the fact that the latter is unordered, and does not have extremity information. This has been implemented for future work that will be described in the next section.

IV. MOVE SEQUENCE PREDICTION

We now move on to the main goal of this project: predicting the move sequence from the holds sequence. This would amount to implementing a systematic solution to bouldering's underlying intellectual challenge of finding the right method (also called *beta*) for a given climb. As we will see, this turned out harder than expected. The general idea was to consider the holds and move sequences as sentences, and use text models to translate one into the other.

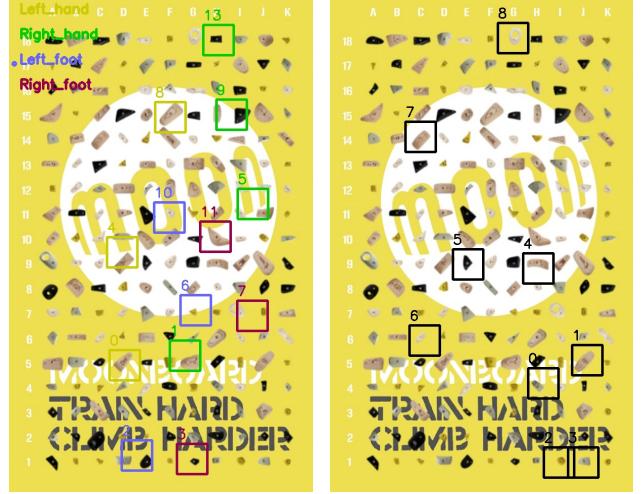


Figure 3: Selection interface with examples of move sequence (left) and holds sequence (right)

A. First tests: adaptation of a Sequence to Sequence network

We first looked into text translation using sequence to sequence (*seq2seq*) translation models. We refer to [7] for the basic structure of the network. It is a PyTorch implementation, relying on an *Encoder* and a *Decoder*, combined with an *Attention* mechanism. The data is encoded into a 512-dimensional latent space, before being processed and mapped onto the output.

The dataset consists of videos used in the students projects, i.e competitions videos provided by the Swiss Olympic Climbing team. The move sequences are extracted from [1], and the holds sequence are computed by clustering the latter with DBSCAN. Regarding our *seq2seq* model, the input sequence is the holds sequence, processed to form a sentence of words of the form " $x_N _ y_N$ ", while the output sequence is the move sequence, processed in the form of a sentence of words such as " $limb_x_N _ y_N$ ".

The model is trained using *Teacher forcing*, and predicts one word at a time from the previous ones. We report in Figure 4 the evolution of the loss with the training epochs.

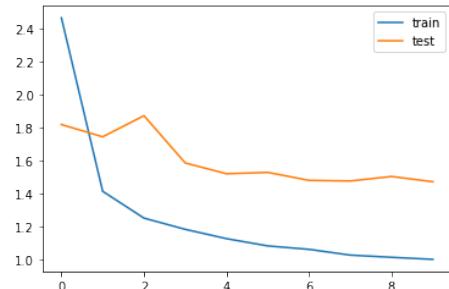


Figure 4: *Seq2seq* model loss trend as a function of the epoch

As we can see, the training went well. In fact, in order to really test the model, we trained it for 100 epochs, which took about two hours on Google Colab's GPU. Nonetheless, the results were disappointing. In fact, as can be seen on Figure 5, the model does not predict the move sequence accurately, and most of the predicted positions are not even holds. This can come from multiple factors, one of them being the technique used to create the holds sequence from the move sequence by clustering.

Another issue comes from the dataset itself. While the framework designed in [1] gives decent results, a lot of move sequences are actually imprecise, with some unpredictability. The videos themselves are not very clean, and exhibit a big diversity of moves. This can impede the model from picking up patterns and accurately learn the translation.

Finally, and arguably the most important, this choice of model constrains us to define the input and output vocabulary to consider every possible combination of {limb, x -coordinate, y -coordinate}. Since the vocabulary is discrete, the coordinates themselves must be discretized on a grid by rounding them up to a given number of decimals. Another problem is that Colab's memory is limited, so that we cannot create arbitrarily big vocabularies. In fact, we were only able to make it work by rounding the coordinates to **one decimal**, i.e. a coordinate can only take the values $\{0, 0.1, 0.2, \dots, 1\}$. This hugely impacts the performance of the model, since it can only predict moves on this grid, which can be far from the actual position. This can in particular explain why the prediction in Figure 5 is so different from the truth, the predicted positions not being accurate.

It is worth mentioning that using the visualization pipeline described in subsection III-A, the movements of the generated skeleton are not that far off the real climber's, keeping in mind that the holds are not correctly placed.

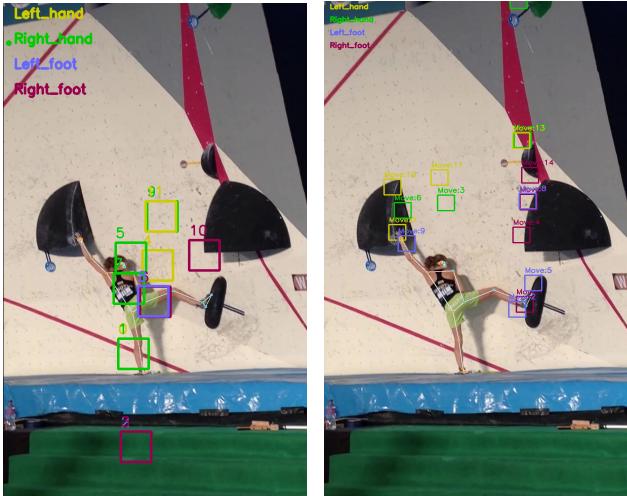


Figure 5: *Seq2seq* move sequence prediction (left) against the truth (right)

B. Next try: Transformer model with Position Embedding

V. CONCLUSIONS

REFERENCES

- [1] C. Beauville, R. Debalme, and T. Patron, “Move sequence detection on bouldering problems.”
- [2] L. Comeliau, J. Heiniger, and W. Wang, “Understanding bouldering using ml methods.”
- [3] C. L. et al., “Mediapipe: A framework for building perception pipelines.”
- [4] M. Ester, H.-P. Kriegel, J. Sander, and X. Xu., “A density-based algorithm for discovering clusters in large spatial databases with noise.”
- [5] “Opencv: Open source computer vision library.” [Online]. Available: <https://github.com/opencv/opencv>
- [6] “Scikit-learn: Machine learning in python.” [Online]. Available: <https://scikit-learn.org/stable/>
- [7] “Nlp from scratch: translation with a sequence to sequence network and attention.” [Online]. Available: https://pytorch.org/tutorials/intermediate/seq2seq_translation_tutorial.html